



○ Run 1 ◇ Run 2


(a) Complete service time




(b) Delivered output rate




(c) Mean request flow




(d) Returned output



(e) Total expert H2D traffic



(f) Correct requests



Natural EOS; 16 requests/run. Bars: available-run means; points: individual complete runs (ordered within arm). Run-level n is shown per panel; small samples, no CI. Drained time includes final drain and excludes initialization/warmup. Actual output counts differ; times are not equal-work speedups. OLMoE; equal 5.5 GiB expert + GPU KV pools (artificial budget; excludes other allocations). Host chunk gaps are not token-level ITL.