


New-host service outcomes · separate run groups

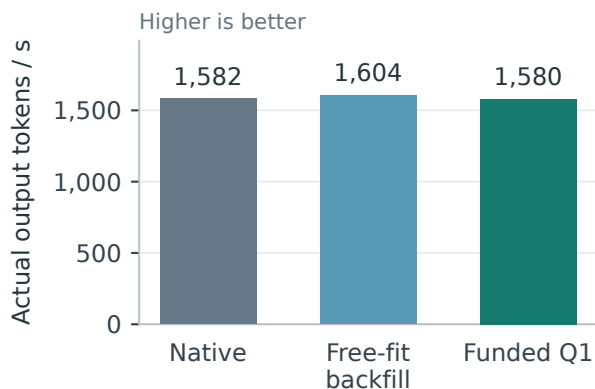
OLMoE-1B-7B (base) · article workload · 128/128 requests complete in every arm · RTX 5090

Strong baselines

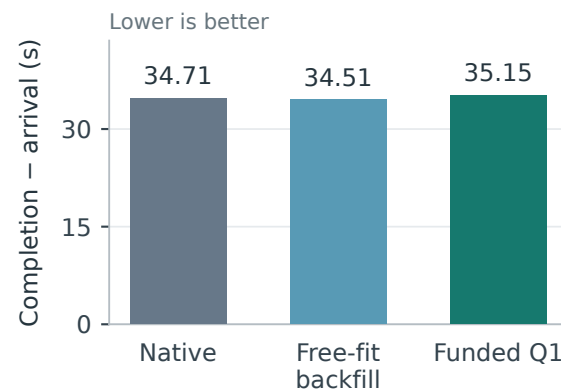
(a) Maximum-gap tail



(b) Output throughput

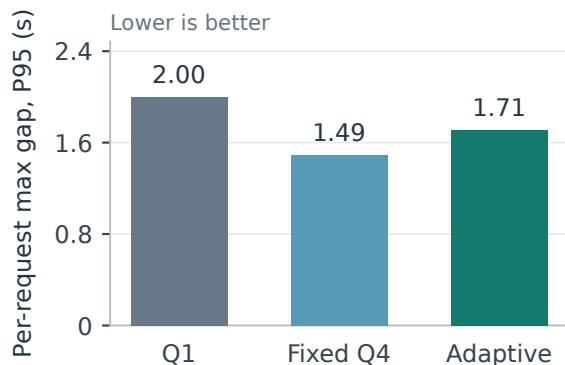


(c) Mean request flow

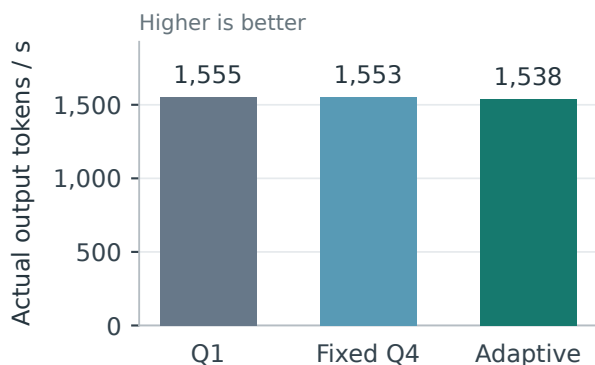


Lease comparison (separate run group)

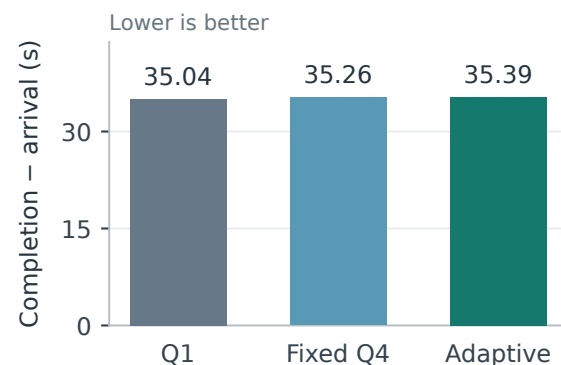
(d) Maximum-gap tail



(e) Output throughput



(f) Mean request flow



Lease block: adaptive is worse than fixed Q4 on all three displayed metrics.

Exploratory: one run per arm; no confidence intervals or pooling across machines.

Rows are separate execution groups. Independent trajectories; actual output counts differ.