


Lease ranking changes across the two run blocks

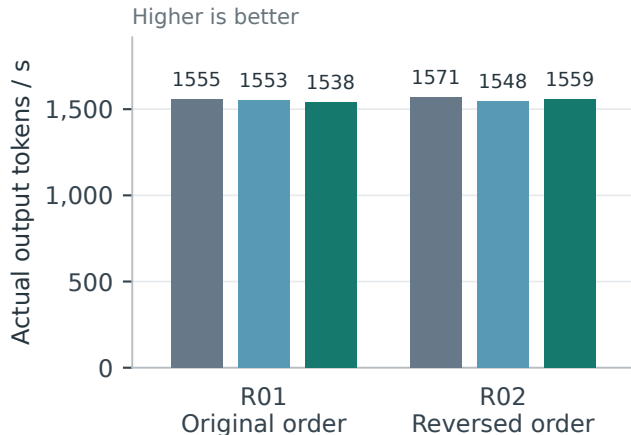
Same base-model article inputs · 128/128 complete in every arm · one run per arm per block · RTX 5090

■ Q1 ■ Fixed Q4 ■ Adaptive

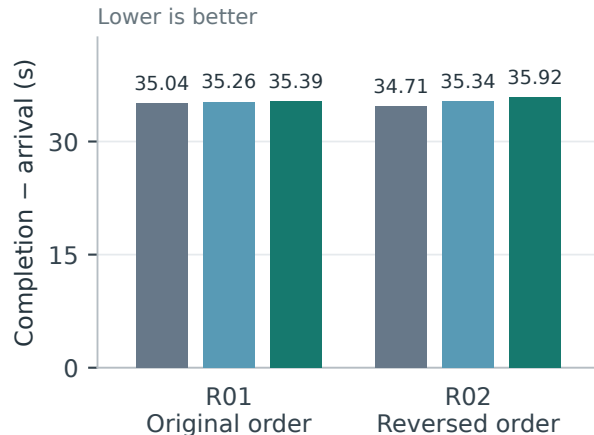
(a) Maximum-gap tail



(b) Output throughput



(c) Mean request flow



Adaptive vs fixed Q4: P95 gap +14.4% → -10.8%; output rate -1.0% → +0.7%; mean flow +0.35% → +1.62%.

R01 order: Q1 → fixed Q4 → adaptive. R02 order: adaptive → fixed Q4 → Q1. Each block is shown separately.

No pooled estimate or confidence intervals. Output trajectories differ; two blocks do not distinguish ordering effects from run noise.