


# Host-observed request timelines

Native cap32: preemption allowed

Safe cap29: guard enabled



One matched repeat; descriptive. Both panels use the same time scale.  
 All spans use host-observed token receipt times. The generation span includes scheduling, recomputation and delivery gaps; it is not a measurement of pure GPU execution time.