Qwen3 MoE expert latency across four formats

COMPLETE: 56 measured points

d=2048, intermediate=768, E=128, top-k=8 | RTX PRO 6000 SM120 | CUDA graph replay

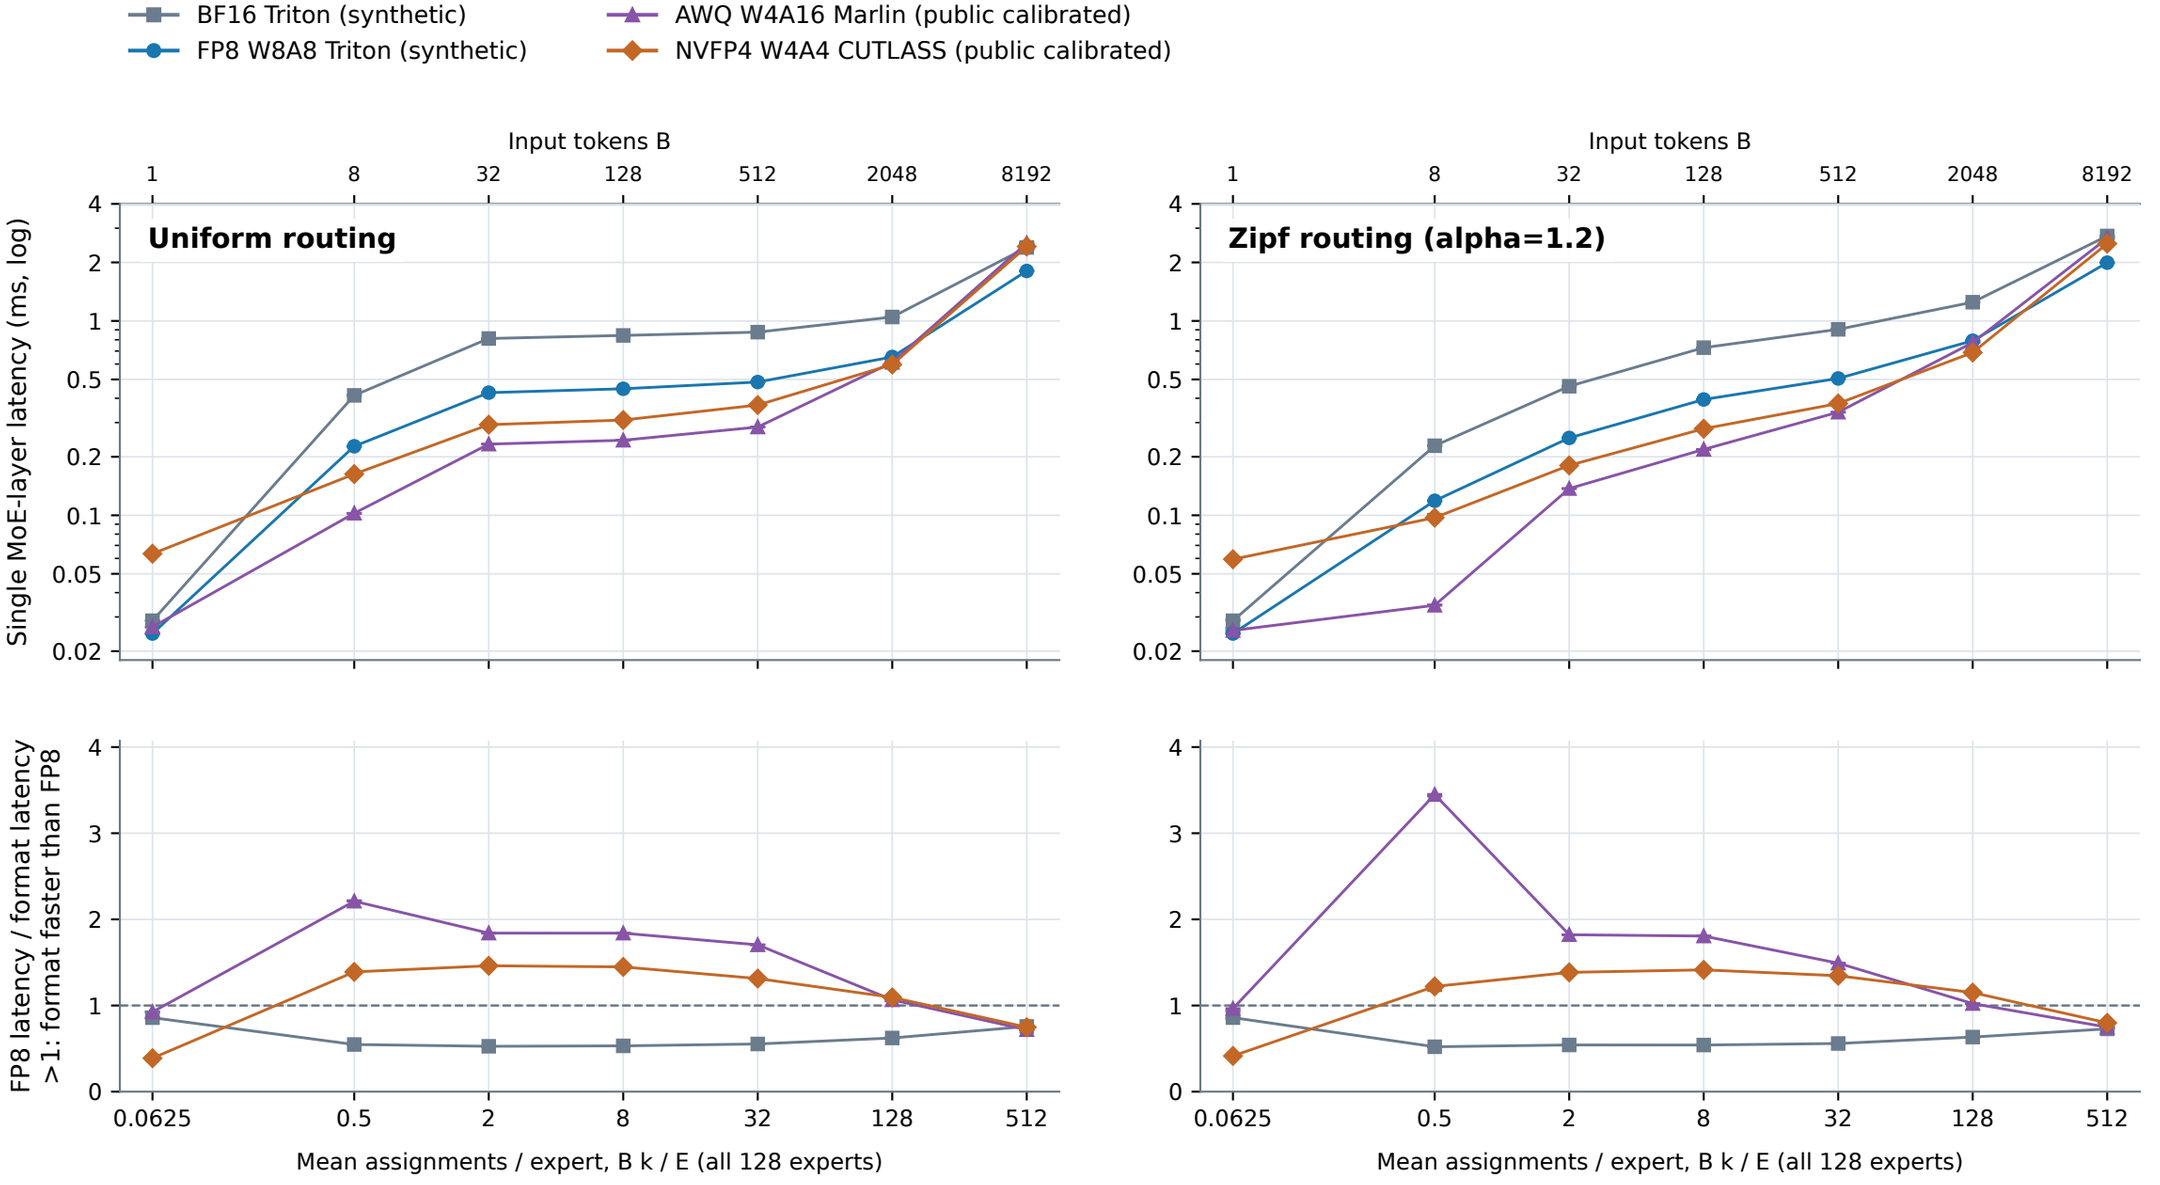

Different weight sources: synthetic BF16/FP8; dtometzki AWQ and NVIDIA NVFP4 calibrated layer 0.

Medians and min-max of five within-process repetitions. Ratio whiskers are unpaired range envelopes, not confidence intervals.

AWQ numerical guard: 3 non-bitwise points pass recorded BF16-roundoff bounds; not exact equality or task-quality validation.

Lines join observed points only; crossings are bracketed, not fitted. All four formats observed at all 14 routing/load points.