Qwen3 expert latency: three measured formats

d = 2048, intermediate = 768, E = 128, top-k = 8 | RTX PRO 6000 SM120 | CUDA graph replay

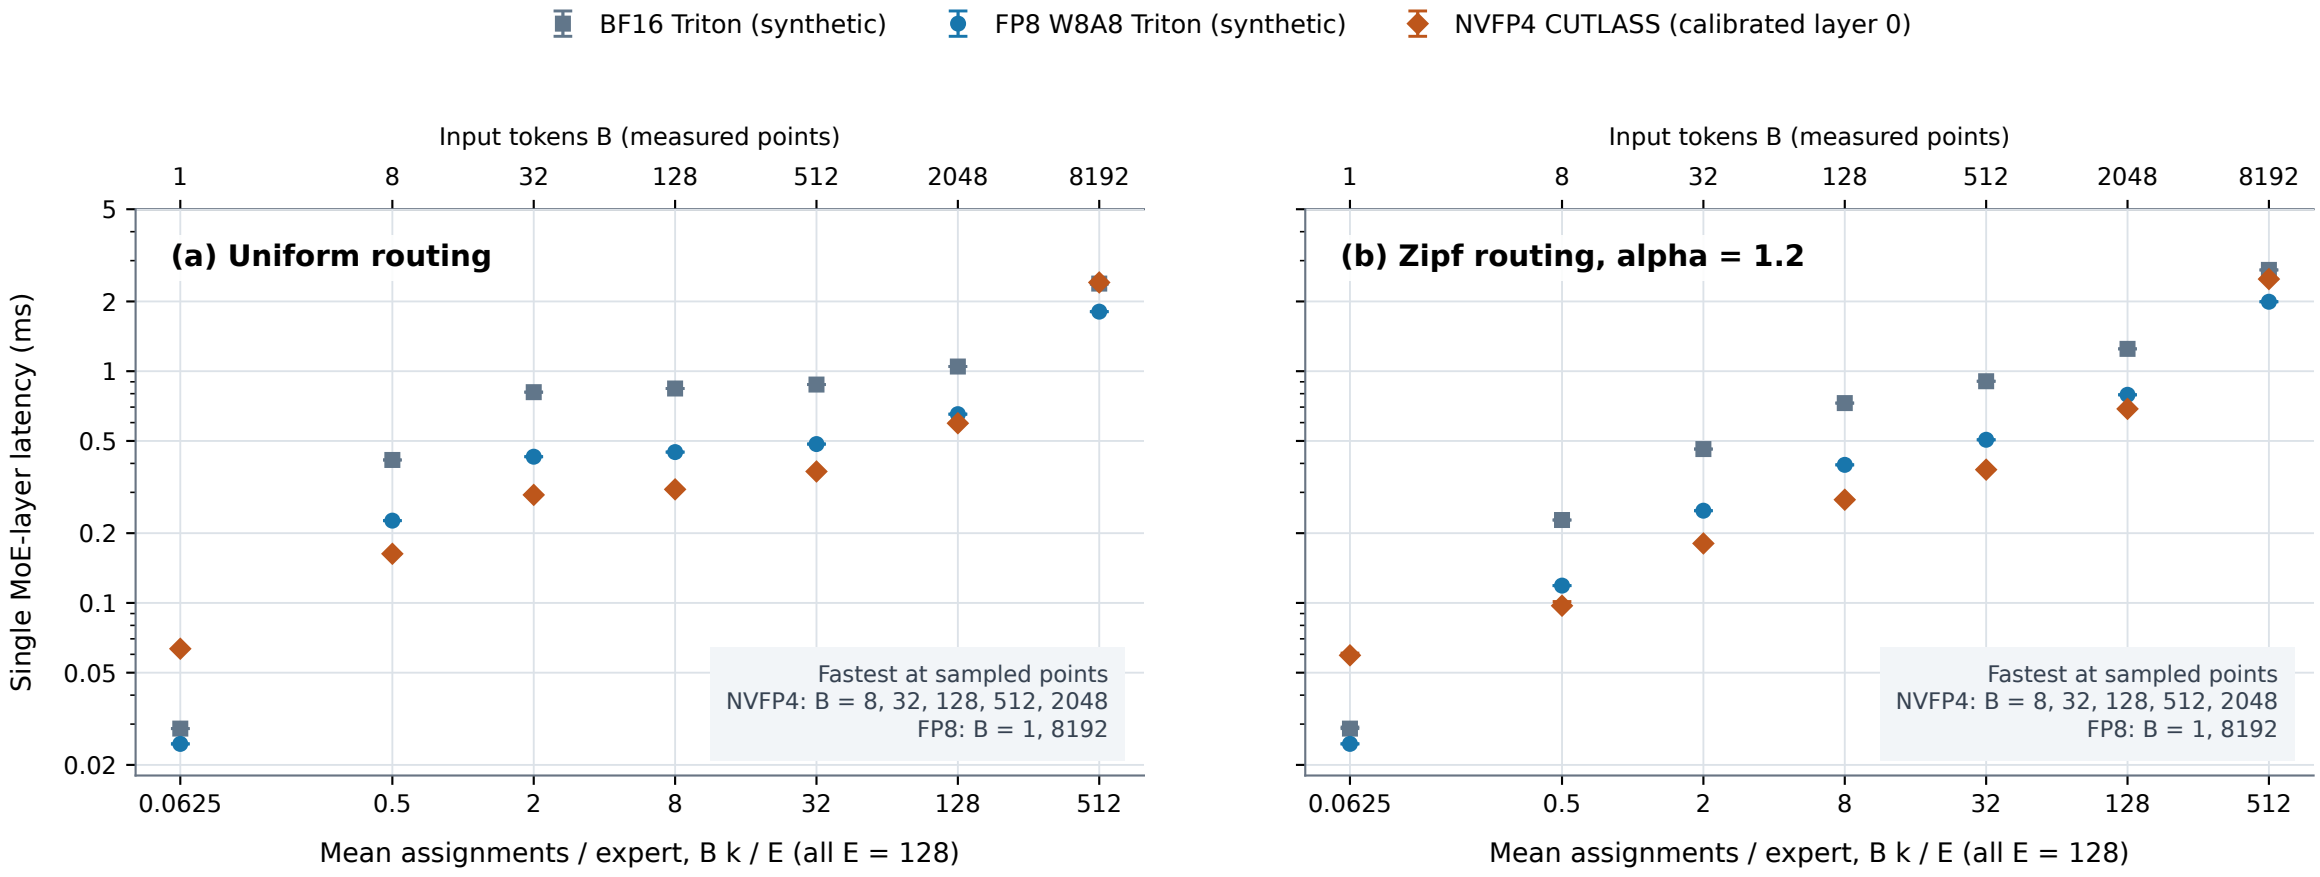

Different weight sources: this is not a same-weight quality or serving comparison.

Markers: medians; whiskers: min-max of five repeats within one process/residency (not independent runs).

Measured B = 1-8192 only. Rank changes are bracketed by B = [1, 8] and [2048, 8192]; no crossover is fitted.

Synthetic routing is identical across formats at each point. Calibrated NVFP4: NVIDIA Qwen3 layer 0. AWQ remains unmeasured.