F2 | Device energy per input token for one MoE layer

CUDA graph replay; Synthetic resident weights; constructed uniform / Zipf routes

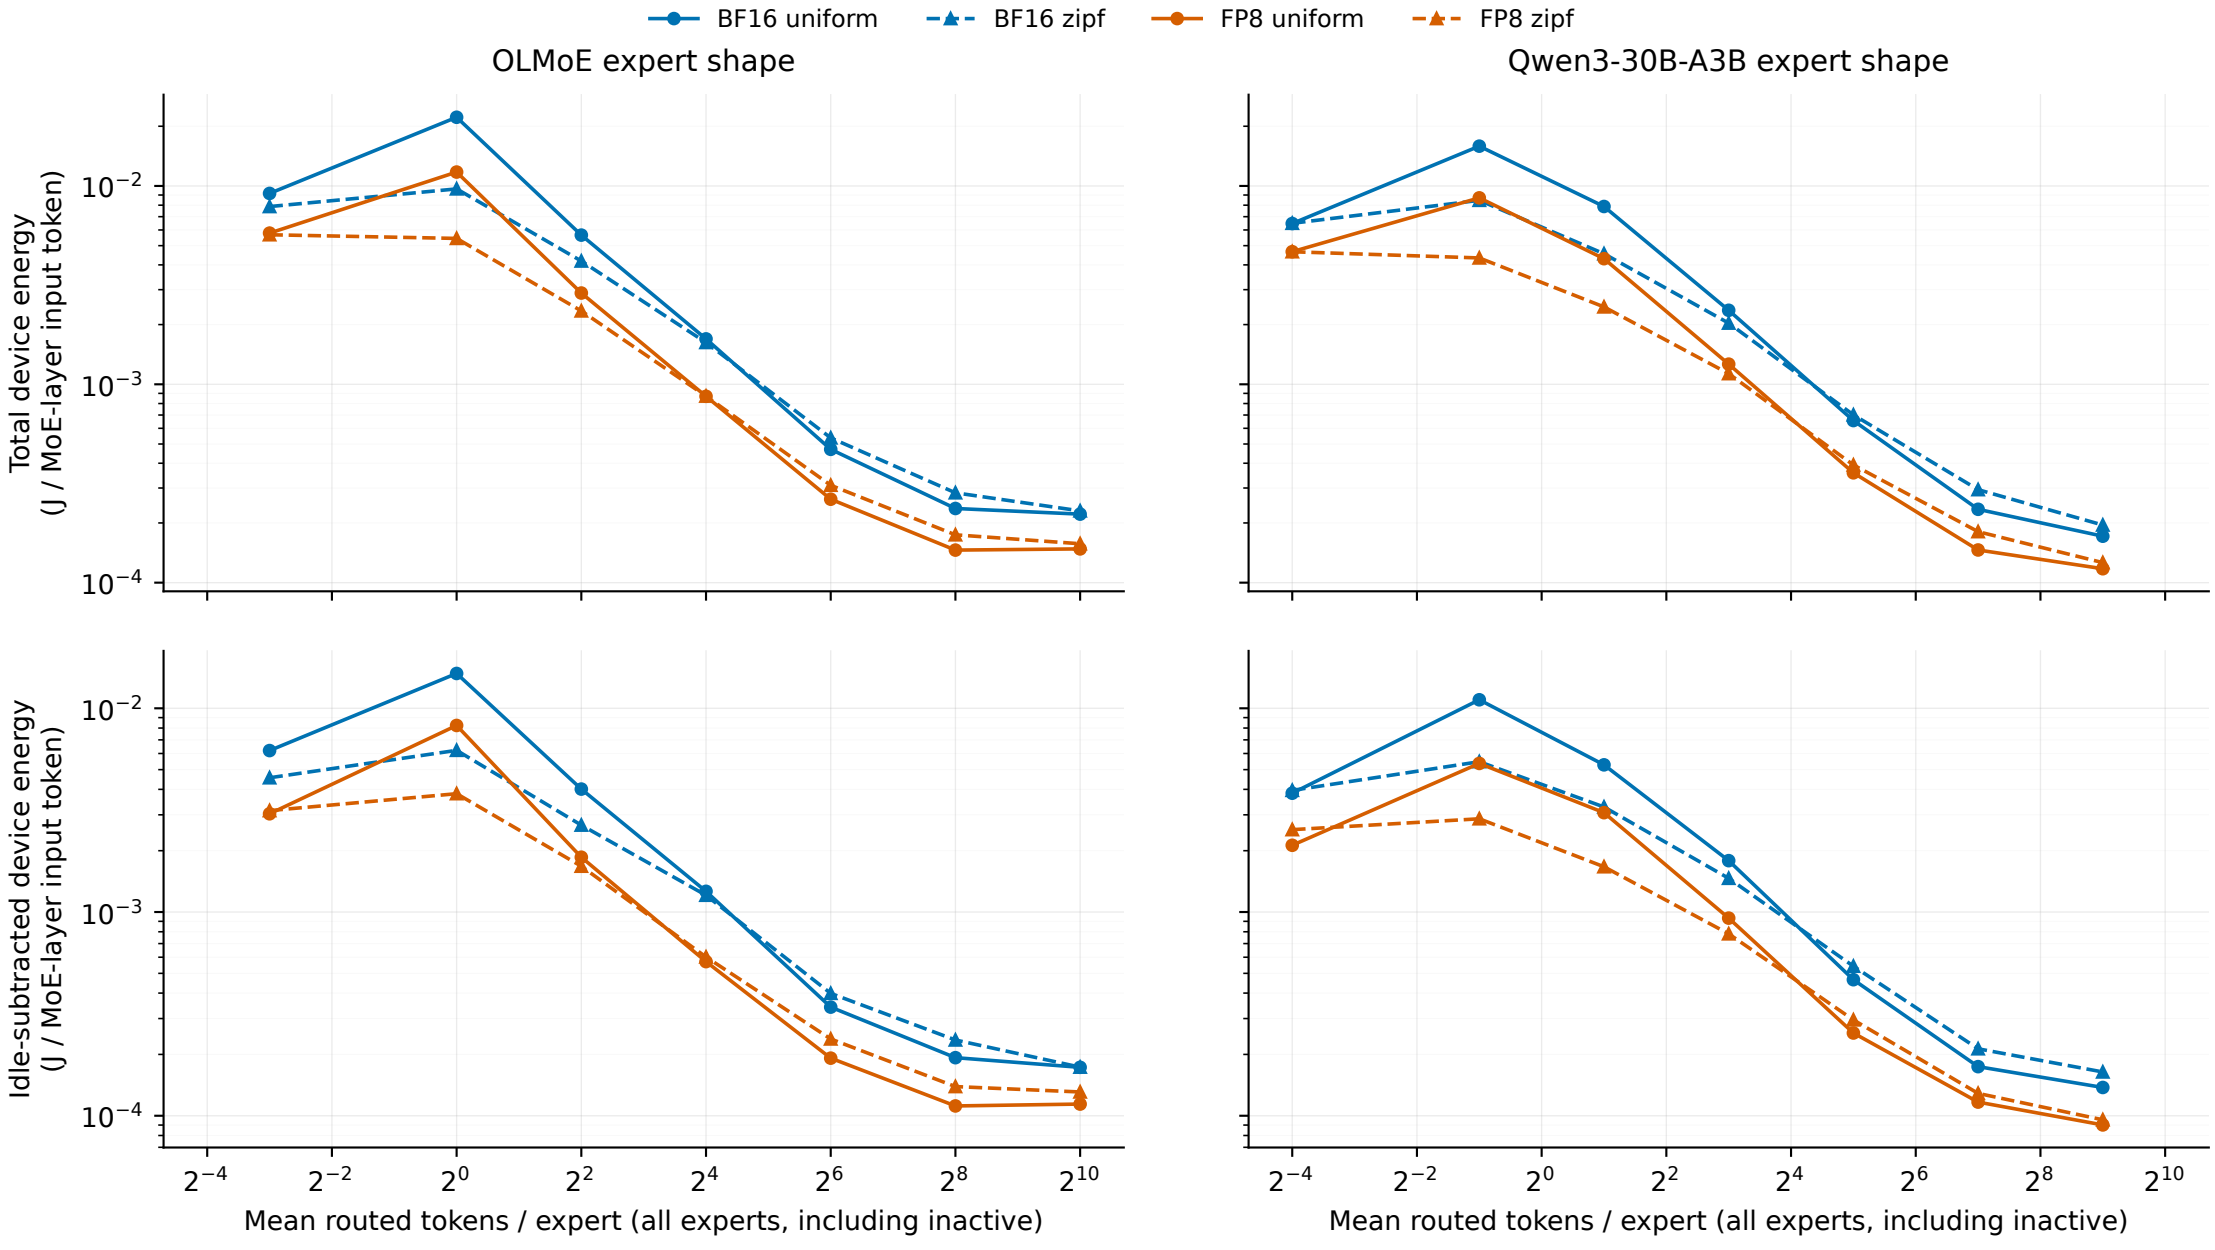

Each marker uses one measured power window / kernel calls / input tokens; energy has no repeat error bars.
These are single-layer input tokens, not generated tokens or serving J/token. Missing points break lines.
Primary energy: nvml_total_energy_counter. Counter vs sampled-integral differences remain in summary.csv.
Coverage: 56/56 planned cases.