F5 Measured batch-dependent benefit | ONE_MODEL_PARTIAL

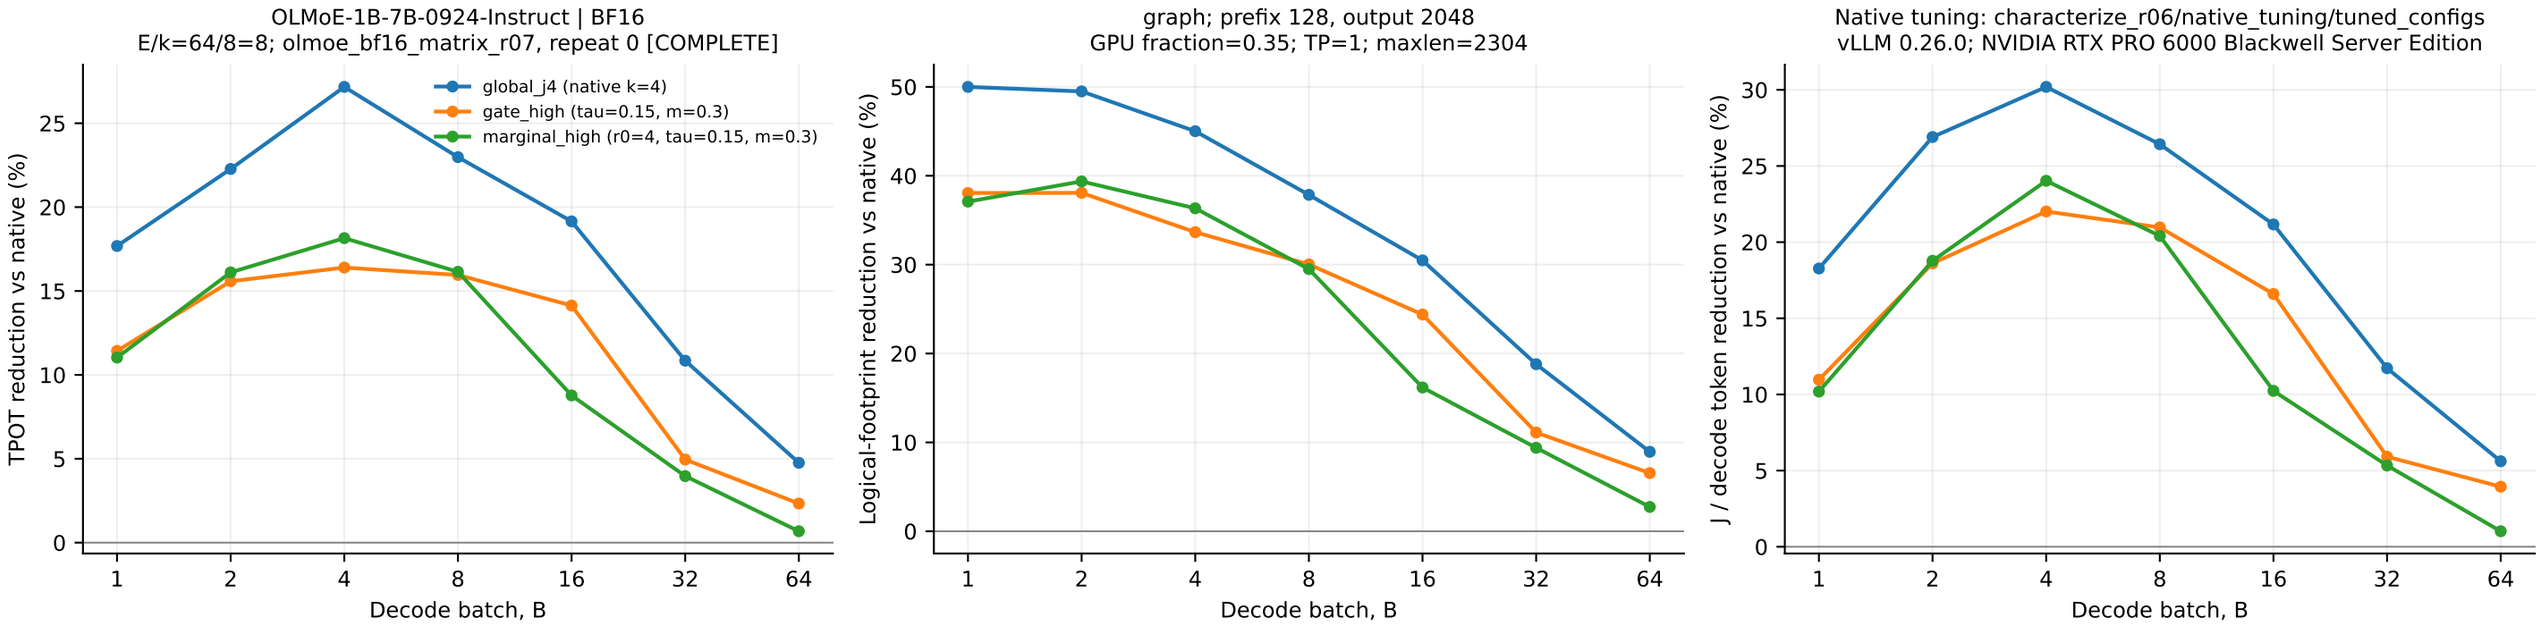

Each row is one run/repeat/condition; formats, tuning, resource budgets and runs are not pooled. Points are measured batches, not independent token repetitions.
E/k is descriptive: models also differ in architecture, routing, KV layout and workload behavior; this figure does not isolate E/k causally.
Footprint compares realized policy trajectories with native, not within-route deletion or hardware DRAM traffic. Energy uses board-power integration; hollow points have a <3 s or unknown window in either paired arm.
Zero is a descriptive baseline, not a significance or deployment threshold. Missing models/formats have no curve; no interpolation of a supported benefit region.