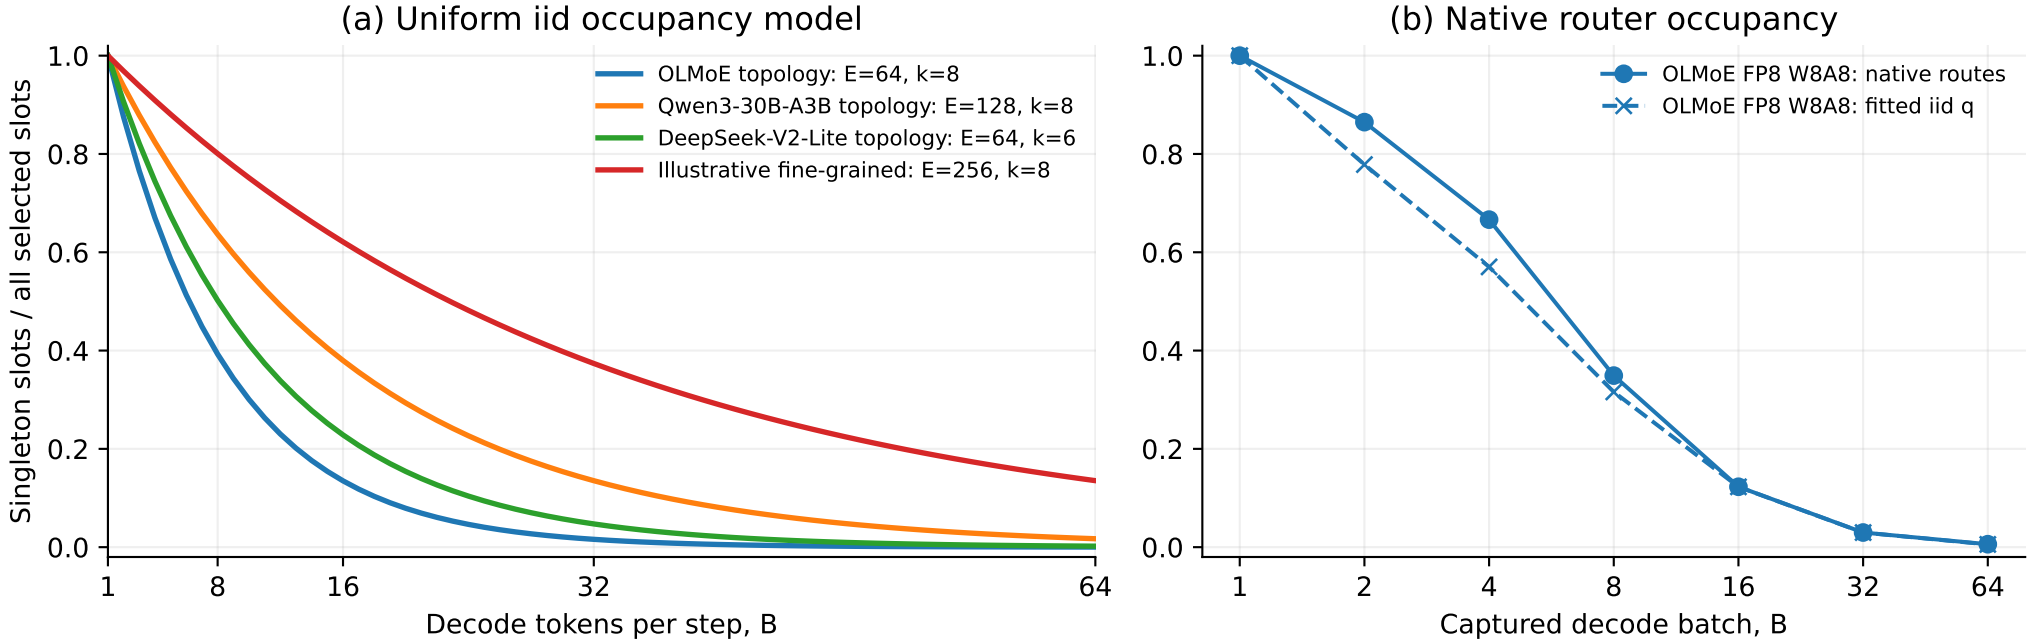


OLMoE FP8 W8A8: 128-token prefixes, 16 decode steps x 16 layers per B; one short capture per batch.
Native router counts measure occupancy; fitted q is in-sample and assumes iid tokens.
Occupancy alone does not measure DRAM bytes or serving latency.