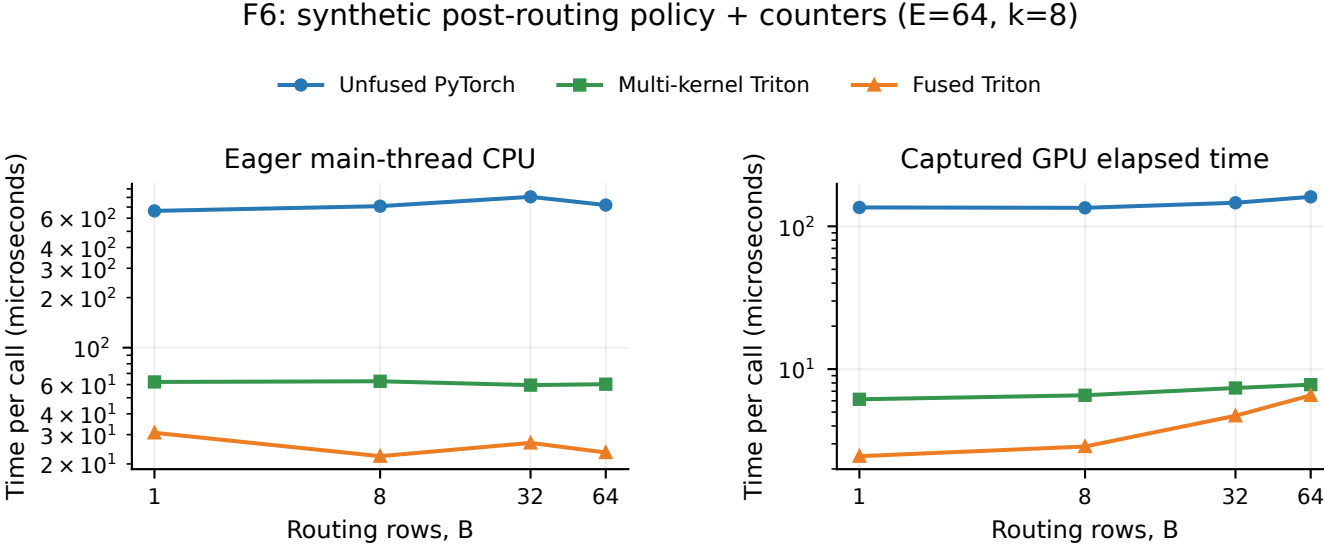
F6: synthetic post-routing policy + counters (E=64, k=8)
Unfused PyTorch
Multi-kernel Triton
Fused Triton
Eager main-thread CPU
Captured GPU elapsed time
Time per call (microseconds)
Routing rows, B
1
8
32
64
One source execution; separate CPU/device clock boundaries, not a subtraction.
CPU excludes final explicit CUDA synchronization, off-CPU waiting, and other threads.