F2 SHORT DIAGNOSTIC | model | bfloat16 | quant=none | triton | eager | NVIDIA RTX PRO 6000 Blackwell Server Edition

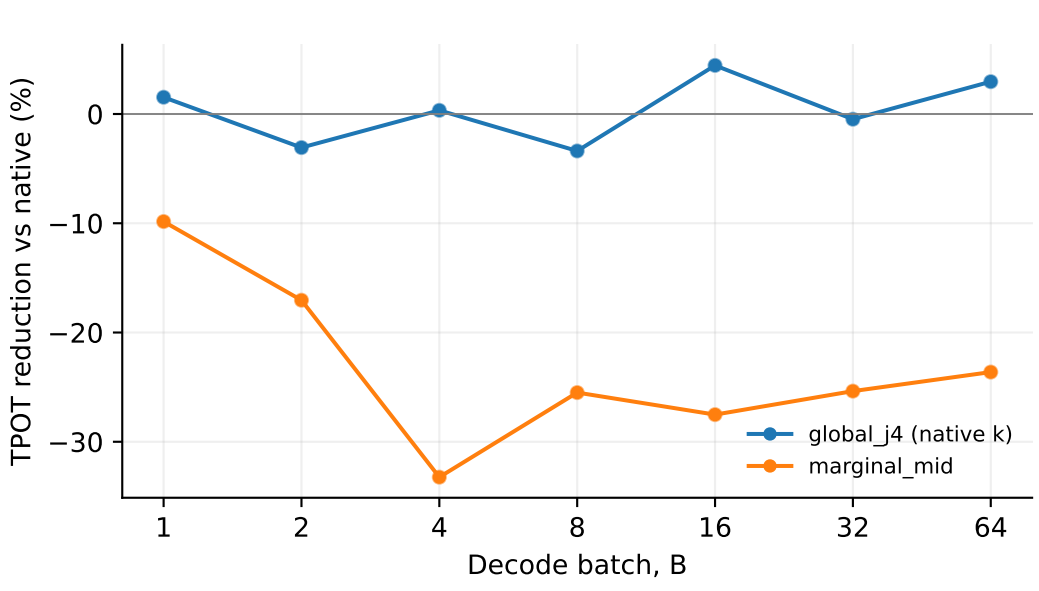

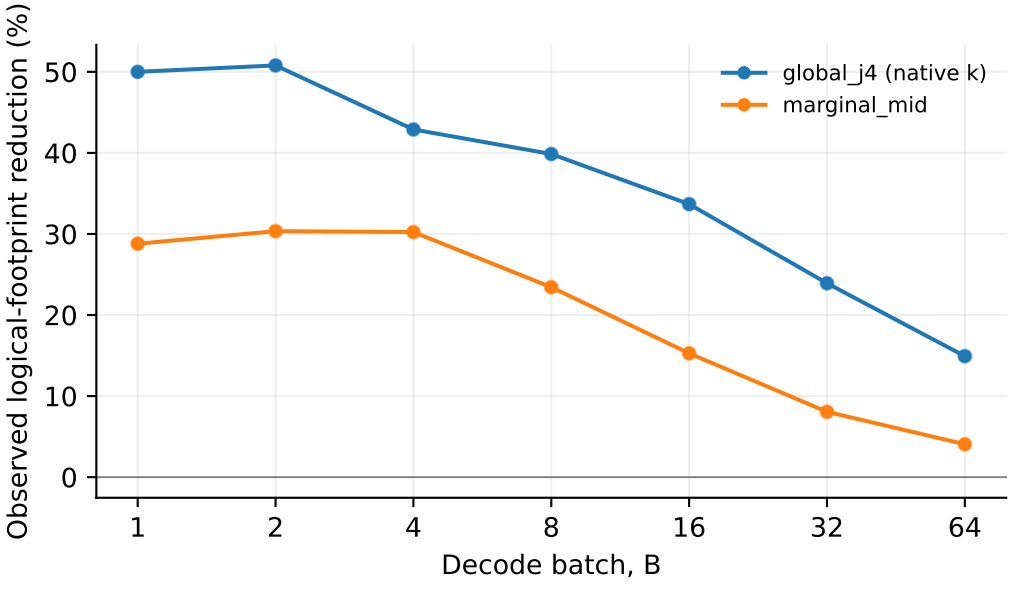

One run per arm / batch; 32-token prompts, 8 fixed outputs (7 decode steps). No uncertainty estimate.
Matched native baseline uses default, untuned MoE configuration. Marginal timing includes policy and counters.
Logical footprint counts distinct expert weight bytes on each arm trajectory; it is not hardware DRAM traffic.
Eager execution, fixed output length; no natural-generation quality or steady-state performance claim.